R2DBC - Reactive Relational Database Connectivity

Ben Hale<bhale@pivotal.io>, Mark Paluch <mpaluch@vmware.com>, Greg Turnquist <gturnquist@pivotal.io>, Jay Bryant <jbryant@pivotal.io>, Elena Felder

Table of Contents

© 2017-2021 The original authors.

Preface

License

Foreword

R2DBC brings a reactive programming API to relational data stores. The Introduction contains more details about its origins and explains its goals.

This document describes the first and initial generation of R2DBC.

Organization of this document

This document is organized into the following parts:

- Introduction
- Goals
- Compliance
- Overview
- Connections
- Transactions
- Statements
- Batches
- Results
- Column and Row Metadata
- Exceptions
- Data Types
- Extensions

Chapter 1. Introduction

R2DBC stands for Reactive Relational Database Connectivity. R2DBC started as an experiment and proof of concept to enable integration of SQL databases into systems that use reactive programming models –- Reactive in the sense of an event-driven, non-blocking, and functional programming model that does not make assumptions over concurrency or asynchronicity. Instead, it assumes that scheduling and parallelization happen as part of runtime scheduling.

1.1. The R2DBC SPI

The R2DBC SPI provides reactive programmatic access to relational databases from Java and other JVM-based programming languages.

R2DBC specifies a service-provider interface (SPI) that is intended to be implemented by driver vendors and used by client libraries. By using the R2DBC SPI, applications written in a JVM programming language can run SQL statements and retrieve results by using an underlying data source. You can also use the R2DBC SPI to interact with multiple data sources in a distributed, heterogeneous environment. R2DBC targets primarily, but is not limited to, relational databases. It aims for a range of data sources whose query and statement interface is based on SQL (or an SQL-like dialect) and that represent their data in a tabular form.

A key difference between R2DBC and imperative data access SPIs is the deferred nature of execution. R2DBC is, therefore, based on Reactive Streams and uses the concepts of `Publisher` and `Subscriber` to allow non-blocking back-pressure-aware data access.

1.2. Requirements and Conventions

R2DBC requires Java 8 as baseline and Reactive Streams in terms of a reactive inter-op API.

A note about modules: R2DBC allow for deployment to JDK 9's module path ("Jigsaw"). For use in Jigsaw-enabled applications, the R2DBC SPI comes with "Automatic-Module-Name" manifest entries which define stable language-level module names ("r2dbc.spi") instead using the artifact name. Of course, R2DBC jars keep working fine on the classpath on both JDK 8 and 9+.

R2DBC's SPI uses non-null semantics by default. Nullable API members are annotated with `@io.r2dbc.spi.Nullable`. This annotation and `@NonNullApi` are meta-annotated with JSR 305 annotations (a dormant but wide-spread JSR). JSR-305 meta-annotations let tooling vendors like IDEA or Kotlin provide null-safety support in a generic way, without having to hard-code support for R2DBC annotations. It is not necessary nor recommended to add a JSR-305 dependency to the project classpath to take advantage of a null-safe API.

R2DBC generally uses zero-based indexes for parameter binding and columns.

1.3. Target Audience

This specification is targeted primarily towards:

- Vendors of drivers that implement the R2DBC SPI.

- Vendors of client implementations who wish to implement a client on top of the R2DBC SPI.
- Vendors of runtime libraries who wish to embed R2DBC into their eco-system to provide R2DBC runtime services.

This specification is also intended to offer:

- An introduction for end-users whose applications use the R2DBC SPI.
- A starting point for developers of other SPIs layered on top of the R2DBC SPI.

Code examples in this specification assume external subscription to `Publisher` objects for brevity.

1.4. Acknowledgements

The R2DBC specification work is being conducted as an effort of individuals that recognized the demand for a reactive, standardized API for relational database access. We want to thank all contributing members for their countless hours of work and discussion.

Thanks also go to Ollie Drotbohm, without whom this initiative would not even exist.

1.5. Following Development

For information on R2DBC source code repositories, nightly builds, and snapshot artifacts, see the R2DBC homepage. You can help make R2DBC best serve the needs of the community by interacting with developers through the community. To follow developer activity, look for the mailing list information on the R2DBC homepage. If you encounter a bug or want to suggest an improvement, please create a ticket on the R2DBC issue tracker. R2DBC has an open-source organization on GitHub that bundles the various projects (SPI and drivers) under R2DBC.

To stay up to date with the latest news and announcements in the R2DBC eco system, you can subscribe to the mailing list. You can also follow the project team on Twitter (@R2DBC).

1.6. Project Metadata

- Version control: https://github.com/r2dbc/r2dbc-spi
- Mailing list: https://groups.google.com/forum/#!forum/r2dbc
- Issue tracker: https://github.com/r2dbc/r2dbc-spi/issues
- Release repository: https://repo1.maven.org/maven2
- Snapshot repository: https://oss.sonatype.org/content/repositories/snapshots

Chapter 2. Goals

This section outlines the goals for R2DBC and the design philosophy for its SPI. It covers the following topics:

- Enabling Reactive Relational Database Connectivity

- Fitting into Reactive JVM platforms

- Offering Vendor-neutral Access to Standard Features

- Embracing Vendor-specific Features

- Keeping the Focus on SQL

- Keeping It Minimal and Simple

- Providing a Foundation for Tools and Higher-level APIs

- Specifying Requirements Unambiguously

2.1. Enabling Reactive Relational Database Connectivity

The R2DBC specification aims for establishing an interface that has a minimal API surface to integrate with relational databases by using a reactive programming model. The most significant goals are honoring and embracing the properties of reactive programming, including the following:

- Non-blocking I/O

- Deferred execution

- Treating application control as a series of events (data, errors, completion, and so on)

- No longer assuming control of resources but leaving resource scheduling to the runtime or platform ("React to resource availability")

- Efficiently using resources

- Leaving flow control to be handled by the runtime

- Stream-oriented data consumption

- Functional programming within operators

- Removing assumptions over concurrency from the programming model and leaving this aspect up the runtime

- Using back-pressure to allow flow control, deferring the actual execution and not overwhelming consumers

2.2. Fitting into Reactive JVM platforms

R2DBC aims for seamless integration of reactive JVM platforms, targeting Java as its primary platform. R2DBC is intended to also be usable from other platforms (such as Kotlin or Scala) without scarifying its SPI for the sake of idiomatic use on a different platform.

2.3. Offering Vendor-neutral Access to Standard Features

R2DBC SPI strives to provide access to features that are commonly found across different vendor implementations. The goal here is to provide a balance between features that are implemented in a driver and these that are better implemented in a client library.

2.4. Embracing Vendor-specific Features

Each database comes with its own feature set and how these are implemented. R2DBC's goal is to define a minimal standard over commonly used functionality and allow for vendor-specific deviation. Drivers can implement additional functionality or make these transparent through the R2DBC SPI.

2.5. Keeping the Focus on SQL

The focus of R2DBC is on accessing relational data from the Java programming language by using databases that provide an SQL interface with which to interact.

The goal here is not to limit implementations to relational-only databases. Instead, the goal is to provide guidance for uniform reactive data access by using tabular data consumption patterns.

2.6. Keeping It Minimal and Simple

R2DBC does not aim to be a general-purpose data-access API.

R2DBC specializes in reactive data access and common usage patterns that result from relational data interaction. R2DBC does not aim to abstract common functionality that needs to be re-implemented by driver vendors in a similar manner. It aims to leave this functionality to client libraries, of which there are typically fewer implementations than drivers.

2.7. Providing a Foundation for Tools and Higher-level APIs

The R2DBC SPI aims for being primarily consumed though client library implementations.

It does not aim to be an end-user or application developer programming interface.

Having a uniform reactive relational data access SPI makes R2DBC a valuable target platform for tool vendors and application developers who want to create portable tools and applications.

2.8. Specifying Requirements Unambiguously

The requirements for R2DBC compliance aim to be unambiguous and easy to identify. The R2DBC specification and the API documentation (Javadoc) clarify which features are required and which are optional.

Chapter 3. Overview of Changes

The following section reflects the history of changes.

3.1. 0.9

Extended transaction definitions

- Introduction of the `TransactionDefinition` interface.

- Introduction of the `Connection.beginTransaction(TransactionDefinition)` method.

Improved Bind Parameter declararation

- Support to set IN, IN/OUT, and OUT parameters.

- Introduction of the `Parameter` interface type hierarchy and the `Parameters` class providing factory methods.

Consumption of OUT Parameters

- Introduction of OUT Parameters.

- Introduction of the `OutParameters`, `OutParametersMetadata`, and `OutParameterMetadata` interfaces.

- Introduction of the `Readable` and `ReadableMetadata` interfaces and retrofitting of `Row` and `ColumnMetadata` interfaces.

- Introduction of the `Result.map(Function<Readable, T>)` method.

Consumption of Result Segments

- Introduction of the `Result.Segment` interface type hierarchy and `Result.filter(Predicate<Segment>)` and `Result.flatMap(Function<Segment, Publisher<T>>)` methods.

Lifecycle extension

- Introduction of the `Lifecycle` interface.

Refinement of `Option`

- Removal of generic type of `ConnectionFactoryOptions.getValue(Option)` and `ConnectionFactoryOptions.getRequiredValue(Option)`.

Refinement of `RowMetadata`

- Deprecate `RowMetadata.getColumnNames()` and introduce `RowMetadata.contains(String)` to simplify constraints and usage around column presence checks.

Lock Wait and Statement Timeouts

- Introduction of Limiting the time for lock acquisition and Limiting execution time for statements.

- Introduction of the `Connection.setLockWaitTimeout(Duration)` and `Connection.setStatementTimeout(Duration)` methods.

3.2. 0.8

Initial version.

Chapter 4. Compliance

This chapter identifies the required features of a D2DBC driver implementation to claim compliance. Any features not identified here are considered optional.

4.1. Definitions

To avoid ambiguity, we will use the following terms in the compliance section and across this specification:

R2DBC driver implementation

(short form: **driver**)A driver that implements the R2DBC SPI. A driver may provide support for features that are not implemented by the underlying database or expose functionality that is not declared by the R2DBC SPI (See **Extension**).

Supported feature

A feature for which the R2DBC SPI implementation supports standard syntax and semantics.

Partially supported feature

A feature for which some methods are implemented with standard syntax and semantics and some required methods are not implemented (typically covered by `default` interface methods).

Extension

A feature that is not covered by R2DBC or a non-standard implementation of a feature that is covered.

Fully implemented

Term to express that an interface has all its methods implemented to support the semantics defined in this specification.

Must implement

Term to express that an interface must be implemented, although some methods on the interface are considered optional. Methods that are not implemented rely on the `default` implementation.

4.2. Guidelines and Requirements

The following guidelines apply to R2DBC compliance:

- An R2DBC SPI should implement SQL support as its primary interface. R2DBC does not rely upon (nor does it presume) a specific SQL version. SQL and aspects of statements can be entirely handled in the data source or as part of the driver.

- The specification consists of this specification document and the specifications documented in each interface's Javadoc.

- Drivers supporting parametrized statements must support bind parameter markers.

- Drivers supporting parametrized statements must support at least one parameter binding method (indexed or named).

- Drivers must support transactions.

- Index references to columns and parameters are zero-based. That is, the first index begins with `0`.

4.3. R2DBC SPI Compliance

A driver that is compliant with the R2DBC specification must do the following:

- Adhere to the guidelines and requirements listed under Guidelines and Requirements.

- Support `ConnectionFactory` discovery through Java Service Loader of `ConnectionFactoryProvider`.

- Implement a non-blocking I/O layer.

- Fully implement the following interfaces:

    - `io.r2dbc.spi.ConnectionFactory`

    - `io.r2dbc.spi.ConnectionFactoryMetadata`

    - `io.r2dbc.spi.ConnectionFactoryProvider`

    - `io.r2dbc.spi.Result`

    - `io.r2dbc.spi.Row`

    - `io.r2dbc.spi.RowMetadata`

    - `io.r2dbc.spi.Batch`

- Implement the `io.r2dbc.spi.Connection` interface, except for the following optional methods:

    - `createSavepoint(⋯)`: Calling this method should throw an `UnsupportedOperations` exception for drivers that do not support savepoints.

    - `releaseSavepoint(⋯)`: Calling this method should be a no-op for drivers that do not support savepoint release.

    - `rollbackTransactionToSavepoint(⋯)`: Calling this method should throw an `UnsupportedOperations` exception for drivers that do not support savepoints.

- Implement the `io.r2dbc.spi.Statement` interface, except for the following optional methods:

    - `returnGeneratedValues(⋯)`: Calling this method should be a no-op for drivers that do not support key generation.

    - `fetchSize(⋯)`: Calling this method should be a no-op for drivers that do not support fetch size hints.

- Implement the `io.r2dbc.spi.ColumnMetadata` interface, except for the following optional methods:

    - `getPrecision()`

    - `getScale()`

    - `getNullability()`

    - `getJavaType()`

    - `getNativeTypeMetadata()`

A driver can implement optional Extensions if it is able to provide extension functionality specified by R2DBC.

Chapter 5. Overview

R2DBC provides an API for Java programs to access one or more sources of data. In the majority of cases, the data source is a relational DBMS and its data is accessed using SQL. R2DBC drivers are not limited to RDBMS but can be implemented on top of other data sources, including stream-oriented systems and object-oriented systems. A primary motivation for R2DBC SPI is to provide a standard API for reactive applications to integrate with a wide variety of data sources.

This chapter gives an overview of the API and the key concepts of the R2DBC SPI. It includes the following topics:

- Establishing a Connection

- Using `ConnectionFactory` Discovery

- R2DBC Connection URL

- Running SQL and Retrieving Results

5.1. Establishing a Connection

R2DBC uses the `Connection` interface to define a logical connection API to the underlying data source. The structure of a connection depends on the actual requirements of a data source and how the driver implements these.

In a typical scenario, an application that uses R2DBC connects to a target data source byusing one of two mechanisms:

- `ConnectionFactories`: R2DBC SPI provides this fully implemented class. It provides `ConnectionFactory` discovery functionality for applications that want to obtain a connection without using a vendor-specific API. When an application first attempts to connect to a data source, `ConnectionFactories` automatically loads any R2DBC driver found on the classpath by using Java's `ServiceLoader` mechanism. See ConnectionFactory Discovery for the details of how to implement the discovery mechanism for a particular driver.

- `ConnectionFactory`: A `ConnectionFactory` is implemented by a driver and provides access to `Connection` creation. An application that wants to configure vendor-specific aspects of a driver can use the vendor-specific `ConnectionFactory` creation mechanism to configure a `ConnectionFactory`.

5.1.1. Using `ConnectionFactory` Discovery

As mentioned earlier, R2DBC supports the concept of discovery to find an appropriate driver for a connection request. Providing a `ConnectionFactory` to an application is typically a configuration infrastructure task. Applications that wish to bootstrap an R2DBC client typically handle this aspect directly in application code and, so, discovery can become a task for application developers.

`ConnectionFactories` provides two standard mechanisms to bootstrap a `ConnectionFactory`:

- URL-based: R2DBC supports a uniform URL-based configuration scheme with a well-defined structure and well-known configuration properties. URLs are represented as Java `String` and

can be passed to `ConnectionFactories` for `ConnectionFactory` lookup.

- Programmatic: In addition to a URL-based configuration, R2DBC provides a programmatic approach so that applications can supply structured configuration options to obtain a `ConnectionFactory`.

In addition to the two preceding methods, R2DBC embraces a mixed mechanism as typical configuration infrastructure mixes URL- and programmatic-based configuration of data sources for enhanced flexibility. A typical use case is the separation of concerns in which data-source coordinates are supplied by using a URL while login credentials originate from a different configuration source.

5.1.2. R2DBC Connection URL

R2DBC defines a standard URL format that is an enhanced form of [RFC 3986 Uniform Resource Identifier (URI): Generic Syntax](#) and its amendments supported by Java's `java.net.URI` type.

The following listing shows the syntax Components from [RFC3986](#):

```
     URI          = scheme ":" driver [ ":" protocol ] ":" hier-part [ "?" query ] [
"#" fragment ]

     scheme       = "r2dbc" / "r2dbcs"

     driver       = ALPHA *( ALPHA )

     protocol     = ALPHA *( ALPHA / DIGIT / "+" / "-" / "." / ":")

     hier-part    = "//" authority path-abempty
                  / path-absolute
                  / path-rootless
                  / path-empty

     authority    = [ userinfo "@" ] host [ ":" port ] [ "," host [ ":" port ] ]

     userinfo     = *( unreserved / pct-encoded / sub-delims / ":" )

     host         = IP-literal / IPv4address / reg-name

     port         = *DIGIT

     path-abempty  = *( "/" segment )
     path-absolute = "/" [ segment-nz *( "/" segment ) ]
     path-rootless = segment-nz *( "/" segment )
     path-empty    = 0<pchar>

     segment       = *pchar
     segment-nz    = 1*pchar
     segment-nz-nc = 1*( unreserved / pct-encoded / sub-delims / "@" )
                ; non-zero-length segment without any colon ":"

     query        = *( pchar / "/" / "?" )

     fragment     = *( pchar / "/" / "?" )

     pct-encoded = "%" HEXDIG HEXDIG

     pchar        = unreserved / pct-encoded / sub-delims / ":" / "@"

     sub-delims  = "!" / "$" / "&" / "'" / "(" / ")"
                / "*" / "+" / "," / ";" / "="

     unreserved  = ALPHA / DIGIT / "-" / "." / "_" / "~"
```

Example 1. R2DBC Connection URL

```
r2dbc:a-driver:pipes://localhost:3306/my_database?locale=en_US
\___/ _____/ \___/    _____/_____/_____/
  |       |      |            |             |          |
scheme  driver protocol  authority       path       query
```

- `scheme`: Identify that the URL is a valid R2DBC URL. Valid schemes are `r2dbc` and `r2dbcs` (configure SSL usage).

- `driver`: Identifier for a driver. The specification has no authority over driver identifiers.

- `protocol`: Used as optional protocol information to configure a driver-specific protocol. Protocols can be organized hierarchically and are separated by a colon (`:`).

- `authority`: Contains an endpoint and authorization. The authority may contain a single host or a collection of hostnames and port tuples by separating these with a comma (`,`).

- `path`: (optional) Used as an initial schema or database name.

- `query`: (optional) Used to pass additional configuration options in the form of `String` key-value pairs by using the key name as the option name.

- `fragment`: Unused (reserved for future use).

`ConnectionFactoryOptions.parse(String)` parses a R2DBC URL into `ConnectionFactoryOptions` using standard and optional extended options. A R2DBC Connection URL is parsed into the following options (by using `ConnectionFactoryOptions` constants):

The following listing shows an example URL:

Example 2. R2DBC Connection URL

```
r2dbc:a-driver:pipes://hello:world@localhost:3306/my_database?locale=en_US
```

The following table describes the standard options:

Table 1. Parsed Standard Options

Option	URL Part	Value as per Example
`ConnectionFactoryOptions.SSL`	`r2dbc`	Unconfigured.
`ConnectionFactoryOptions.DRIVER`	`driver`	`a-driver`
`ConnectionFactoryOptions.PROTOCOL`	`protocol`	`pipes`
`ConnectionFactoryOptions.USER`	User-part of `authority`	`hello`
`ConnectionFactoryOptions.PASSWORD`	Password-part of `authority`	`world`

Option	URL Part	Value as per Example
`ConnectionFactoryOptions.HOST`	Host-part of `authority`	`localhost`
`ConnectionFactoryOptions.PORT`	Port-part of `authority`	`3306`
`ConnectionFactoryOptions.DATAB ASE`	`path` without the leading `/`	`my_database`

The following table describes the extended options:

Table 2. Parsed Extended Options

Option	URL Part	Value as per Example
`locale`	key-value tuple from `query`	`en_US`

> 🛈 R2DBC defines well-known standard options that are available as runtime constants through `ConnectionFactories`. Additional options identifiers are created through `Option.valueOf(…)`.

> 🛈 Note that Connection URL Parsing cannot access `Option` type information `T` due to Java's type erasure. Options configured by URL parsing are represented as `String` values.

Example 3. Obtaining a `ConnectionFactory` using R2DBC URL

```
ConnectionFactory factory = ConnectionFactories.get("r2dbc:a-
driver:pipes://localhost:3306/my_database?locale=en_US");
```

Example 4. Obtaining a `ConnectionFactory` using `ConnectionFactoryOptions`

```
ConnectionFactoryOptions options = ConnectionFactoryOptions.builder()
    .option(ConnectionFactoryOptions.DRIVER, "a-driver")
    .option(ConnectionFactoryOptions.PROTOCOL, "pipes")
    .option(ConnectionFactoryOptions.HOST, "localhost")
    .option(ConnectionFactoryOptions.PORT, 3306)
    .option(ConnectionFactoryOptions.DATABASE, "my_database")
    .option(Option.valueOf("locale"), "en_US")
    .build();

ConnectionFactory factory = ConnectionFactories.get(options);
```

5.2. Running SQL and Retrieving Results

Once a connection has been established, an application using the R2DBC SPI can execute queries and updates against the connected database. The R2DBC SPI provides a text-based command

interface to the most commonly used features of SQL databases. R2DBC driver implementations may expose additional functionality in a non-standard way.

Applications use methods in the `Connection` interface to specify transaction attributes and create `Statement` or `Batch` objects. These statements are used to execute SQL and retrieve results and allow for binding values to parameter bind markers. The `Result` interface encapsulates the results of an SQL query. Statements may also be batched, allowing an application to submit multiple commands to a database as a single unit of execution.

Chapter 6. Connections

R2DBC uses the `Connection` interface to define a logical connection API to the underlying data source. The structure of a connection depends on the actual requirements of the data source and how the driver implements these.

The data source can be an RDBMS, a stream-oriented data system, or some other source of data with a corresponding R2DBC driver. A single application that uses R2DBC SPI can maintain multiple connections to either a single data source or across multiple data sources. From a R2DBC driver perspective, a `Connection` object represents a single client session. It has associated state information, such as user ID and what transaction semantics are in effect. A `Connection` object is not safe for concurrent state-changing by multiple subscribers. A connection object can be shared across multiple threads that serially run operations by using appropriate synchronization mechanisms.

To obtain a connection, the application can:

- Interact with the `ConnectionFactories` class by working with one or more `ConnectionFactoryProvider` implementations.
- Directly interact with a `ConnectionFactory` implementation.

See Establishing a Connection for more details.

6.1. The `ConnectionFactory` Interface

R2DBC drivers must implement the `ConnectionFactory` interface as a mandatory part of the SPI. Drivers can provide multiple `ConnectionFactory` implementations, depending on the protocol in use or aspects that require the use of a different `ConnectionFactory` implementation. The following listing shows the `ConnectionFactory` interface:

Example 5. `ConnectionFactory` Interface

```
public interface ConnectionFactory {

    Publisher<? extends Connection> create();

    ConnectionFactoryMetadata getMetadata();

}
```

The following rules apply:

- A `ConnectionFactory` represents a resource factory for deferred connection creation. It may create connections by itself, wrap a `ConnectionFactory`, or apply connection pooling on top of a `ConnectionFactory`.
- A `ConnectionFactory` provides metadata about the driver itself through

`ConnectionFactoryMetadata`.

- A `ConnectionFactory` uses deferred initialization and must initiate connection resource allocation after requesting the item (`Subscription.request(1)`) and not upon calling `create` itself.

- Connection creation must emit exactly one `Connection` or an error signal.

- Connection creation must be cancellable (`Subscription.cancel()`). Canceling connection creation must release ("close") the connection and all associated resources.

- A `ConnectionFactory` should expect that it can be wrapped. Wrappers must implement the `Wrapped<ConnectionFactory>` interface and return the underlying `ConnectionFactory` when `Wrapped.unwrap()` gets called.

6.1.1. ConnectionFactory Metadata

`ConnectionFactory` instances are required to expose metadata to identify the driver (`ConnectionFactory`) and its capabilities. Metadata must not require a connection to a data source. The following listing shows the `ConnectionFactoryMetadata` interface:

Example 6. `ConnectionFactoryMetadata` *Interface*

```
public interface ConnectionFactoryMetadata {

    String getName();

}
```

See the R2DBC SPI Specification for more details.

6.2. `ConnectionFactory` Discovery Mechanism

As part of its usage, the `ConnectionFactories` class tries to load any R2DBC driver classes referenced by the `ConnectionFactoryProvider` interface listed in the Java Service Provider manifests that are available on the classpath.

Drivers must include a file called `META-INF/services/io.r2dbc.spi.ConnectionFactoryProvider`. This file contains the name of the R2DBC driver's implementation (or implementations) of `io.r2dbc.spi.ConnectionFactoryProvider`. To ensure that drivers can be loaded by using this mechanism, `io.r2dbc.spi.ConnectionFactoryProvider` implementations are required to provide a no-argument constructor. The following listing shows a typical `META-INF/services/io.r2dbc.spi.ConnectionFactoryProvider` file:

Example 7. META-INF/services/io.r2dbc.spi.ConnectionFactoryProvider file contents

```
com.example.ConnectionFactoryProvider
```

The following listing shows the `ConnectionFactoryProvider` interface:

Example 8. `ConnectionFactoryProvider` *Interface*

```
public interface ConnectionFactoryProvider {

    ConnectionFactory create(ConnectionFactoryOptions connectionFactoryOptions);

    boolean supports(ConnectionFactoryOptions connectionFactoryOptions);

    String getDriver();

}
```

`ConnectionFactories` uses a `ConnectionFactoryOptions` object to look up a matching driver by using a two-step model:

1. Look up an adequate `ConnectionFactoryProvider`.
2. Obtain the `ConnectionFactory` from the `ConnectionFactoryProvider`.

`ConnectionFactoryProvider` implementations are required to return a `boolean` indicator whether or not they support a specific configuration represented by `ConnectionFactoryOptions`. Drivers must expect any plurality of `Option` instances to be configured. Drivers must report that they support a configuration only if the `ConnectionFactoryProvider` can provide a `ConnectionFactory` based on the given `ConnectionFactoryOptions`. A typical task handled by `supports` is checking driver and protocol options. Drivers should gracefully fail if a `ConnectionFactory` creation through `ConnectionFactoryProvider.create(⋯)` is not possible (i.e. when required options were left unconfigured). The `getDriver()` method reports the driver identifier that is associated with the `ConnectionFactoryProvider` implementation to provide diagnostic information to users in misconfiguration cases.

See the R2DBC SPI Specification and ConnectionFactory Discovery for more details.

6.3. The `ConnectionFactoryOptions` Class

The `ConnectionFactoryOptions` class represents a configuration for a request a `ConnectionFactory` from a `ConnectionFactoryProvider`. It enables the programmatic connection creation approach without using driver-specific classes. `ConnectionFactoryOptions` instances are created by using the builder pattern, and properties are configured through `Option<T>` identifiers. A `ConnectionFactoryOptions` is immutable once created. `Option` objects are reused as part of the built-in constant pool. Options are identified by a literal.

`ConnectionFactoryOptions` defines a set of well-known options:

Table 3. Well-known Options

Constant	URL Literal	Type	Description
SSL	ssl	java.lang.Boolean	Whether the connection is configured to require SSL.
DRIVER	driver	java.lang.String	Driver identifier.
PROTOCOL	protocol	java.lang.String	Protocol details, such as the network protocol used to communicate with a server.
USER	user	java.lang.String	User account name.
PASSWORD	password	java.lang.CharSequence	User or database password.
HOST	host	java.lang.String	Database server name.
PORT	port	java.lang.Integer	Database server port number.
DATABASE	database	java.lang.String	Name of the particular database on a server.
CONNECT_TIMEOUT	connectTimeout	java.time.Duration	Connection timeout to obtain a connection.
LOCK_WAIT_TIMEOUT	lockWaitTimeout	java.time.Duration	Lock acquisition timeout.
STATEMENT_TIMEOUT	statementTimeout	java.time.Duration	Statement timeout.

The following rules apply:

- The set of options is extensible.

- Drivers can declare which well-known options they require and which they support.

- Drivers can declare which extended options they require and which they support.

- Drivers should not fail in creating a connection if more options are declared than the driver consumes, as a ConnectionFactory should expect to be wrapped.

- Connection URL Parsing cannot access Option type information T due to Java's type erasure. Options obtained by URL parsing beyond well-known keys are represented as String values.

The following example shows how to set options for a ConnectionFactoryOptions:

Example 9. Configuration of `ConnectionFactoryOptions`

```
ConnectionFactoryOptions options = ConnectionFactoryOptions.builder()
    .option(ConnectionFactoryOptions.HOST, "…")
    .option(Option.valueOf("tenant"), "…")
    .option(Option.sensitiveValueOf("encryptionKey"), "…")
    .build();
```

See the R2DBC SPI Specification for more details.

6.4. Obtaining `Connection` Objects

Once a `ConnectionFactory` is bootstrapped, connections are obtained from the `create()` method. The following example shows how to obtain a connection:

Example 10. Obtaining a `Connection`

```
// factory is a ConnectionFactory object
Publisher<? extends Connection> publisher = factory.create();
```

The connection is active once it has been emitted by the `Publisher` and must be released ("closed") once it is no longer in use.

6.5. Connection Metadata

Connections are required to expose metadata about the database they are connected to. Connection Metadata is typically discovered dynamically based from information obtained during `Connection` initialization.

Example 11. `ConnectionMetadata` *Interface*

```
public interface ConnectionMetadata {

    String getDatabaseProductName();

    String getDatabaseVersion();

}
```

See the R2DBC SPI Specification for more details.

6.6. Validating `Connection` Objects

The `Connection.validate(…)` method indicates whether the `Connection` is still valid. The `ValidationDepth` argument passed to this method indicates the depth to which a connection is validated: `LOCAL` or `REMOTE`.

- `ValidationDepth.LOCAL`: Requests client-side-only validation without engaging a remote conversation to validate a connection.

- `ValidationDepth.REMOTE`: Initiates a remote validation by issuing a query or other means to validate a connection and the remote session.

If `Connection.validate(…)` emits `true`, the `Connection` is still valid. If `Connection.validate(…)` emits `false`, the `Connection` is not valid, and any attempt to perform database interaction fails. Callers of this method do not expect error signals or empty completion.

6.7. Closing `Connection` Objects

Calling `Connection.close()` prepares a close handle to release the connection and its associated resources. Connections must be closed to ensure proper resource disposal. You can use `Connection.validate(…)` to determine whether a `Connection` has been closed or is still valid. The following example shows how to close a connection:

Example 12. Closing a `Connection`

```
// connection is a ConnectionFactory object
Publisher<Void> close = connection.close();
```

See the R2DBC SPI Specification for more details.

6.8. Limiting the time for lock acquisition

The lock acquisition timeout can be configured on the `Connection` level allowing to associate the connection with a lock acquisition limit. The default lock acquisition timeout is vendor-specific and can be specified for new connections through connection factory options.

If the timeout is exceeded, a `R2dbcTimeoutException` is raised to the client. Support for transaction-bound timeouts through `TransactionDefinition` is subject to vendor-specific availability.

6.9. Limiting execution time for statements

The statement timeout can be configured on the `Connection` level allowing to associate the connection with a timeout limit. By default there is no limit on the amount of time allowed for a running statement to complete unless specified through connection factory options. The minimum amount of time can be used to either let the data source cancel the statement or attempt a client-side cancellation.

Once the data source has had an opportunity to process the request to terminate the running command, a `R2dbcTimeoutException` is raised to the client and no additional processing can occur against the previously running command without re-executing a statement.

The timeout applies to statements through a combination of `Statement#execute`, `Batch#execute` and `Result` consumption. In the case of `Batch`/`Statement` batching, it is vendor-specific whether the timeout is applied to individual SQL commands or the entire batch.

Support for transaction-bound timeouts through `TransactionDefinition` is subject to vendor-specific availability.

Chapter 7. Transactions

Transactions are used to provide data integrity, isolation, correct application semantics, and a consistent view of data during concurrent database access. All R2DBC-compliant drivers are required to provide transaction support. Transaction management in the R2DBC SPI reflects SQL concepts:

- Transaction Boundaries

- Auto-commit mode

- Transaction isolation levels

- Savepoints

This section explains transaction semantics associated with a single `Connection` object.

7.1. Transaction Boundaries

Transactions begin implicitly or explicitly. Implicit transactions begin by starting SQL execution when a `Connection` is in auto-commit mode (which is the default for newly created connections). Alternatively, explicit transactions begin by invoking the `beginTransaction()` method that disables auto-commit mode. Transactions are started by either the R2DBC driver or its underlying database.

Newly started transactions inherit attributes from the connection such as Isolation Level. Starting a transaction using `beginTransaction(TransactionDefinition)` allows R2DBC drivers to compute a transaction definition from any plurality of attributes before starting the actual transaction. A driver may apply optimizations such as reduction of database roundtrips. Attributes retrieved from `TransactionDefinition` are only valid during the transaction.

7.1.1. TransactionDefinition Interface

The `io.r2dbc.spi.TransactionDefinition` interface provides a mechanism for drivers to obtain transaction attributes when starting a transaction using a vendor-neutral API. R2DBC drivers query objects implementing the interface for supported attributes. The interface defines attribute constants for commonly supported attributes such as Isolation Level, Transaction Mutability, Transaction Name, and Lock Wait Timeout. Drivers may specify and query additional attributes.

Usage

Starting a transaction using a `TransactionDefinition` object allows R2DBC drivers to use the definition object as callback to query for transactional attributes. The following example shows how to start a transaction:

Example 13. Starting a transaction using `TransactionDefinition`.

```
// connection is a Connection object
// definition is a TransactionDefinition object
Publisher<Void> begin = connection.beginTransaction(definition);
```

> ℹ️ Drivers may reject attributes based on the connection configuration. An example would be the use of read-write transactions with connections in read-only mode.

Interface Methods

The following methods are available on the `TransactionDefinition` interface:

- `getAttribute`

The `getAttribute` **Method**

R2DBC drivers invoke the `getAttribute` method to query for transaction attributes. Attributes that have no configuration value attached are considered unconfigured and therefore represented with a `null` value. The value type follows the actual attribute identified by a `Option` constant.

7.2. Auto-commit Mode

A `ConnectionFactory` creates new `Connection` objects with auto-commit mode enabled, unless specified otherwise through connection configuration options. The `Connection` interface provides two methods to interact with auto-commit mode:

- `setAutoCommit`
- `isAutoCommit`

R2DBC applications should change auto-commit mode by invoking the `setAutoCommit` method instead of executing SQL commands to change the connection configuration. If the value of auto-commit is changed during an active transaction, the current transaction is committed. If `setAutoCommit` is called and the value for auto-commit is not changed from its current value, this is treated as a no-op.

Changing auto-commit mode typically engages database activity. Therefore, the method returns a `Publisher`. Querying auto-commit mode is typically a local operation that involves driver state without database communication.

When auto-commit is disabled, you must explicitly start and clean up each transaction by calling the `Connection` methods `beginTransaction` and `commitTransaction` (or `rollbackTransaction`), respectively.

This is appropriate for cases where transaction management is being done in a layer above the driver, such as:

- The application needs to group multiple SQL statements into a single transaction.

- An application container manages the transaction state.

7.3. Transaction Isolation

Transaction isolation levels define the level of visibility ("isolation") for statements that are run within a transaction. They impact concurrent access while multiple transactions are active.

The default transaction level for a `Connection` object is vendor-specific and determined by the driver that supplied the connection. Typically, it defaults to the transaction level supported by the underlying data source.

The `Connection` interface provides two methods to interact with transaction isolation levels:

- `setTransactionIsolationLevel`

- `getTransactionIsolationLevel`

R2DBC applications should change transaction isolation levels by invoking the `setTransactionIsolationLevel` method instead of running SQL commands to change the connection configuration.

Changing transaction isolation levels typically involves database activity. Therefore, the method returns a `Publisher`. Changing an isolation level during an active transaction results in implementation-specific behavior. Querying transaction isolation levels is typically a local operation that involves driver state without database communication. The return value of the method `getTransactionIsolationLevel` reflects the current isolation level when it actually occurs. `IsolationLevel` is an extensible runtime-constant so drivers may define their own isolation levels. A driver may not support transaction levels. Calling `getTransactionIsolationLevel` results in returning vendor-specific `IsolationLevel` object.

7.3.1. Performance Considerations

When you increase the transaction isolation level, databases typically require more locking and resource overhead to ensure isolation level semantics. This, in turn, lowers the degree of concurrent access that can be supported. As a result, applications may see degraded performance when they use higher transaction isolation levels. For this reason, a transaction manager, whether it is the application itself or part of the application container, should weigh the need for data consistency against the requirements for performance when determining which transaction isolation level is appropriate.

7.4. Savepoints

Savepoints provide a fine-grained control mechanism by marking intermediate points within a transaction. Once a savepoint has been created, a transaction can be rolled back to that savepoint without affecting preceding work.

7.4.1. Working with Savepoints

The `Connection` interface defines methods to interact with savepoints:

- `createSavepoint`
- `releaseSavepoint`
- `rollbackTransactionToSavepoint`

Savepoints are created during an active transaction and are valid only as long as the transaction is active. You can use the `createSavepoint` method to set a savepoint within the current transaction. A transaction is started if `createSavepoint` is invoked and there is no active transaction (switching the connection to disabled auto-commit mode). The `rollbackTransactionToSavepoint` method is used to roll back work to a previous savepoint without rolling back the entire transaction. the following example shows how to roll back a transaction to a savepoint:

Example 14. Rolling back a transaction to a savepoint

```
// connection is a Connection object
Publisher<Void> begin = connection.beginTransaction();

Publisher<? extends Result> insert1 = connection.createStatement("INSERT INTO
books VALUES ('John Doe')").execute();

Publisher<Void> savepoint = connection.createSavepoint("savepoint");

Publisher<? extends Result> insert2 = connection.createStatement("INSERT INTO
books VALUES ('Jane Doe')").execute();

...

Publisher<Void> partialRollback =
connection.rollbackTransactionToSavepoint("savepoint");

...

Publisher<Void> commit = connection.commit();

// publishers are materialized in the order: begin, insert1, savepoint, insert2,
partialRollback, commit
```

Drivers that do not support savepoint creation and rolling back to a savepoint should throw an `UnsupportedOperationException` to indicate these features are not supported.

7.4.2. Releasing a Savepoint

Savepoints allocate resources on the databases, and some vendors may require releasing a savepoint to dispose resources. The `Connection` interface defines the `releaseSavepoint` method to release savepoints that are no longer needed.

Savepoints that were created during a transaction are released and are invalidated when the transaction is committed or when the entire transaction is rolled back. Rolling a transaction back to a savepoint automatically releases it. A rollback also invalidates any other savepoints that were created after the savepoint in question.

Calling `releaseSavepoint` for drivers that do not support savepoint release results in a no-op.

Chapter 8. Statements

This section describes the `Statement` interface. It also describes related topics, including parameterized statement and auto-generated keys.

8.1. The Statement Interface

The `Statement` interface defines methods for running SQL statements. SQL statements may contain parameter bind markers for input parameters.

8.1.1. Creating Statements

`Statement` objects are created by `Connection` objects, as the following example shows:

Example 15. Creating a non-parameterized `Statement`

```
// connection is a Connection object
Statement statement = connection.createStatement("SELECT title FROM books");
```

Each `Connection` object can create multiple `Statement` objects that the program can concurrently run at any time. Resources that are associated with a statement are released as soon as the connection is closed.

8.1.2. Running Statement Objects

`Statement` objects are run by calling the `execute()` method. Depending on the SQL, the resulting `Publisher` may return one or many `Result` objects. A `Statement` is always associated with its `Connection`. Therefore, the connection state affects `Statement` execution at execution time. The following example shows how to run a statement:

Example 16. Running a `Statement`

```
// statement is a Statement object
Publisher<? extends Result> publisher = statement.execute();
```

8.2. Parameterized Statements

The SQL that is used to create a statement can be parameterized by using vendor-specific bind markers. The portability of SQL statements across R2DBC implementations is not a goal.

Parameterized `Statement` objects are created by `Connection` objects in the same manner as non-parameterized `Statements`. See the the following example:

Example 17. Creating three parameterized `Statement` *objects by using vendor-specific parameter bind markers*

```
// connection is a Connection object
Statement statement1 = connection.createStatement("SELECT title FROM books WHERE
author = :author");

Statement statement2 = connection.createStatement("SELECT title FROM books WHERE
author = @P0");

Statement statement3 = connection.createStatement("SELECT title FROM books WHERE
author = $1");
```

Parameter bind markers are identified by the `Statement` object. Parameterized statements may be cached by R2DBC implementations for reuse (for example, for prepared statement execution).

8.2.1. Binding Parameters

The `Statement` interface defines `bind(…)` and `bindNull(…)` methods to provide parameter values for bind marker substitution. A parameter value consists of the actual value that is bound to a parameter and its type. Using scalar values according to [Mapping of Data Types](#) lets the R2DBC infer the actual database type. Using a `Parameter` object allows for more control over the database type definition. Each bind method accepts two arguments. The first is either an ordinal position parameter starting at `0` (zero) or the parameter placeholder representation. The method of parameter binding (positional or by identifier) is vendor-specific, and a driver should document its preferred binding mechanism. The second and any remaining parameters specify the value to be assigned to the parameter. The following example shows how to bind parameters to a statement object by using placeholders:

Example 18. Binding parameters to a `Statement` *object by using placeholders*

```
// connection is a Connection object
Statement statement = connection.createStatement("SELECT title FROM books WHERE
author = $1 and publisher = $2");
statement.bind("$1", "John Doe");
statement.bind("$2", "Happy Books LLC");
```

Alternatively, parameters can be bound by index, as the following example shows:

Example 19. Binding parameters to a Statement object by index

```
// connection is a Connection object
Statement statement = connection.createStatement("SELECT title FROM books WHERE
author = $1 and publisher = $2");
statement.bind(0, "John Doe");
statement.bind(1, "Happy Books LLC");
```

Binding parameters using a Parameter with a Type allows for more control over the actual database type. io.r2dbc.spi.R2dbcType defines commonly used types.

Example 20. Binding parameters to a Statement object using a Parameter object

```
// connection is a Connection object
Statement statement = connection.createStatement("SELECT title FROM books WHERE
author = $1 and publisher = $2");
statement.bind(0, Parameters.in(R2dbcType.NVARCHAR, "John Doe"));
statement.bind(1, Parameters.in(R2dbcType.VARCHAR, "Happy Books LLC"));
```

A value must be provided for each bind marker in the Statement object before the statement can be run. The execute method validates a parameterized Statement and throws an IllegalStateException if a bind marker is left without a binding.

8.2.2. Batching

Parameterized Statement objects accept multiple parameter binding sets to submit a batch of commands to the database for running. A batch run is initiated by invoking the add() method on the Statement object after providing all parameters. After calling add(), the next set of parameter bindings is provided by calling bind methods accordingly. The following example shows how to run a batch Statement:

Example 21. Running a Statement batch

```
// connection is a Connection object
Statement statement = connection.createStatement("INSERT INTO books (author,
publisher) VALUES ($1, $2)");
statement.bind(0, "John Doe").bind(1, "Happy Books LLC").add();
statement.bind(0, "Jane Doe").bind(1, "Scary Books Inc");
Publisher<? extends Result> publisher = statement.execute();
```

A batch run emits one or many Result objects, depending on how the implementation executes the batch.

8.2.3. Setting `NULL` Parameters

You can use the `bindNull` method to set any parameter to `NULL`. It takes two parameters:

- Either the ordinal position of the bind marker or the name.
- The value type of the parameter.

The following example shows how to set `NULL` value:

Example 22. Setting a `NULL` value using type inference.

```
// statement is a Statement object
statement.bindNull(0, String.class);
```

Typed `Parameter` objects representing a null value can be bound either by calling `bindNull(⋯)` or `bind(⋯)`:

Example 23. Setting a typed `NULL` value.

```
// statement is a Statement object
statement.bind(0, Parameters.in(R2dbcType.VARCHAR));
```

> 🛈 | Not all databases allow for a non-typed `NULL` to be sent to the underlying database.

8.2.4. Setting `IN/OUT` and `OUT` Parameters

Statements can take three kinds of parameters:

- `IN` parameters (default type) as described in Binding Parameters.
- `IN/OUT` parameters.
- `OUT` parameters.

The parameter can be specified as either an ordinal parameter or a named parameter. A value must be set for each parameter marker in the statement that represents an `IN` or `IN/OUT` parameter. `OUT` parameters are generally not associated with a value and may require a type hint.

Parameter types can either make use of type inference by specifying the value or a Java `Class` or reference a `Type`.

`IN/OUT` parameters are assigned values whose result can be retrieved after running the statement as described in OUT Parameters. Parameters are assigned using the `bind(⋯)` methods as described in Binding Parameters.

The following example shows how to set a `IN/OUT` parameter:

Example 24. Setting a `IN/OUT` parameter value using type inference.

```
// connection is a Connection object
Statement statement = connection.createStatement("CALL my_proc ($1)");
statement.bind(0, Parameters.inOut("John Doe"));
```

`OUT` parameters are value-less parameters whose result can be retrieved after running the statement as described in OUT Parameters. Parameters are assigned using the `bind(…)` methods as described in Binding Parameters.

The following example shows how to set a `OUT` parameter:

Example 25. Setting a `OUT` parameter value using type inference.

```
// connection is a Connection object
Statement statement = connection.createStatement("CALL my_proc ($1)");
statement.bind(0, Parameters.out(String.class));
```

8.3. Retrieving Auto Generated Values

Many database systems provide a mechanism that automatically generates a value when a row is inserted. The value that is generated may or may not be unique or represent a key value, depending on the SQL and the table definition. You can call the `returnGeneratedValues` method to retrieve the generated value. It tells the `Statement` object to retrieve generated values. The method accepts a variable-argument parameter to specify the column names for which to return generated keys. The emitted `Result` exposes a column for each automatically generated value (taking the column name hint into account). The following example shows how to retrieve auto-generated values:

Example 26. Retrieving auto-generated values

```
// connection is a Connection object
Statement statement = connection.createStatement("INSERT INTO books (author,
publisher) VALUES ('John Doe', 'Happy Books LLC')").returnGeneratedValues("id");
Publisher<? extends Result> publisher = statement.execute();

// later
result.map((readable) -> readable.get("id"));
```

When column names are not specified, the R2DBC driver implementation determines the columns or value to return.

See the R2DBC SPI Specification for more details.

8.4. Performance Hints

The `Statement` interface provides a method that you can use to provide hints to a R2DBC driver. Calling `fetchSize` applies a fetch-size hint to each query produced by the statement. Hints provided to the driver through this interface may be ignored by the driver if they are not appropriate or supported.

Back-pressure hints can be used by drivers to derive an appropriate fetch size. To optimize for performance, it can be useful to provide hints to the driver on a per-statement basis to avoid unwanted interference of back-pressure hint propagation.

Note that back-pressure should be considered a utility for flow control and not to limit the result size. Result size limitations should be part of the query statement.

Chapter 9. Batches

This section describes the `Batch` interface.

9.1. The Batch Interface

The `Batch` interface defines methods for running groups of SQL statements. SQL statements may not contain parameter bind markers for input parameters. A batch is created to run multiple SQL statements for performance reasons.

9.1.1. Creating Batches

`Batch` objects are created by `Connection` objects, as the following example shows:

Example 27. Creating a `Batch`

```
// connection is a Connection object
Batch batch = connection.createBatch();
```

Each `Connection` object can create multiple `Batch` objects that can be used concurrently by the program and can be run at any time. Resources that are associated with a batch are released as soon as the connection is closed.

9.1.2. Executing Batch Objects

`Batch` objects are run by calling the `execute()` method after adding one or more SQL statements to a `Batch`. The resulting `Publisher` returns a `Result` object for each statement in the batch. A `Batch` is always associated with its `Connection`. Therefore, the connection state affects `Batch` execution at run time.

The following example shows how to run a batch:

Example 28. Running a `Batch`

```
// connection is a Connection object
Batch batch = connection.createBatch();
Publisher<? extends Result> publisher = batch.add("SELECT title, author FROM
books")
    .add("INSERT INTO books VALUES('John Doe', 'HappyBooks LLC')")
    .execute();
```

See the R2DBC SPI Specification for more details.

Chapter 10. Results

This section explains the `Result` interface and the related `Readable` interface. It also describes related topics, including result consumption.

10.1. Result Characteristics

`Result` objects are forward-only and read-only objects that allow consumption of the following result types:

- Update count

- Tabular results

- OUT parameters

Results move forward from the first `Row` to the last one. After emitting the last row, drivers can expose out parameters by emitting `OutParameters`. After that, a `Result` object gets invalidated and results from the same `Result` object can no longer be consumed. Rows and parameters contained in the result depend on how the underlying database materializes the results. That is, it contains the rows that satisfy the query at either the time the query is run or as the rows are retrieved. An R2DBC driver can obtain a `Result` either directly or by using cursors.

`Result` reports the number of rows affected for SQL statements, such as updates for SQL Data Manipulation Language (DML) statements. The update count can be empty for statements that do not modify rows. After emitting the update count, a `Result` object gets invalidated and rows from the same `Result` object can no longer be consumed. The following example shows how to get a count of updated rows:

Example 29. Consuming an update count

```
// result is a Result object
Publisher<Integer> rowsUpdated = result.getRowsUpdated();
```

`Result` represens a stream of result segments. Due to its nature, a result allows consumption of either tabular results, out parameters, or an update count through `map(⋯)` respective `getRowsUpdated(⋯)` but not both. Depending on how the underlying database materializes results, an R2DBC driver can lift this limitation.

A `Result` object is emitted for each statement result in a forward-only direction. A statement can lead to multiple `Result` objects caused by either multiple bindings or by running a statement that materializes multiple result sets.

10.2. Creating Result Objects

A Result object is created as the result of running a Statement object. The Statement.execute() method returns a Publisher that emits Result objects as the result of running the statement. The following example shows how to create a Result object:

Example 30. Creating a Result object

```
// connection is a Connection object
Statement statement = connection.createStatement("SELECT title, author FROM books");
Publisher<? extends Result> results = statement.execute();
```

The Result object emits a Row object for each row in the books table (which contains two columns: title and author). The following sections detail how these rows and columns can be consumed.

10.2.1. Cursor Movement

Result objects can be backed by direct results (that is, a query that returns results directly) or by cursors. By consuming Row objects, an R2DBC driver advances the cursor position. Thus, external cursor navigation is not possible.

Canceling subscription of tabular results stops cursor reads and releases any resources associated with the Result object.

10.3. Rows

A `Row` object represents a single row of tabular results.

10.3.1. Retrieving Values

The `Result` interface provides a `map(⋯)` method for retrieving values from `Row` objects. The `map` method accepts a `BiFunction` (also referred to as mapping function) object that accepts `Row` and `RowMetadata`. The mapping function is called upon row emission with `Row` and `RowMetadata` objects. A `Row` is only valid during the mapping function callback and is invalid outside of the mapping function callback. Thus, `Row` objects must be entirely consumed by the mapping function. The overloaded `map` method accepting a `Function` object for `Readable` can be called with either `Row` or `OutParameters` objects or vendor-specific extensions.

The Column and Row Metadata section contains additional details on metadata.

10.4. Interface Methods

The following methods are available on the `Row` interface:

- `Object get(int)` (inherited from `Readable`)
- `Object get(String)` (inherited from `Readable`)
- `<T> T get(int, Class<T>)` (inherited from `Readable`)
- `<T> T get(String, Class<T>)` (inherited from `Readable`)
- `RowMetadata getMetadata()`

`get(int[, Class])` methods accept column indexes starting at `0`, `get(String[, Class])` methods accept column name aliases as they are represented in the result. Column names used as input to the `get` methods are case-insensitive. Column names do not necessarily reflect the column names as they are in the underlying tables but, rather, how columns are represented (for example, aliased) in the result. The following example shows how to create and consume a `Row` by using its index:

Example 31. Creating and Consuming a `Row` by obtaining a column by index

```
// result is a Result object
Publisher<Object> values = result.map((row, rowMetadata) -> row.get(0));
```

The following example shows how to create and consume a `Row` by using its column name:

Example 32. Creating and Consuming a `Row` by obtaining a column by name

```
// result is a Result object
Publisher<Object> titles = result.map((row, rowMetadata) -> row.get("title"));
```

Calling `get` without specifying a target type returns a suitable value representation according to Mapping of Data Types. When specifying a target type, the R2DBC driver tries to convert the value to the target type. The following example shows how to create and consume a `Row` with type conversion:

Example 33. Creating and Consuming a `Row` with type conversion

```
// result is a Result object
Publisher<String> values = result.map((row, rowMetadata) -> row.get(0,
String.class));
```

You can also consume multiple columns from a `Row`, as the following example shows:

Example 34. Consuming multiple columns from a `Row`

```
// result is a Result object
Publisher<Book> values = result.map((row, rowMetadata) -> {
    String title = row.get("title", String.class);
    String author = row.get("author", String.class);

    return new Book(title, author);
});
```

When the column value in the database is SQL `NULL`, it can be returned to the Java application as `null`.

> ℹ️ `null` values cannot be returned as Reactive Streams values and must be wrapped for subsequent usage.

> ℹ️ Invalidating a `Row` does **not** release `Blob` and `Clob` objects that were obtained from the `Row`. These objects remain valid for at least the duration of the transaction in which they were created, unless their `discard()` method is called.

10.5. OUT Parameters

A `OutParameters` object represents a set of `OUT` parameters as result of a stored procedure/server-side function invocation.

10.5.1. Retrieving Values

The `Result` interface provides a `map(…)` method for retrieving values from `Readable` objects. The `map` method accepts a `Function` (also referred to as mapping function) for `Readable` which can be an implementation of `OutParameters` or `Row` objects or vendor-specific extensions. The mapping function is called upon emission with `Readable` objects. A `Readable` is only valid during the mapping

function callback and is invalid outside of the mapping function callback. Thus, `Readable` objects must be entirely consumed by the mapping function.

The OUT Parameter Metadata section contains additional details on metadata.

10.6. Interface Methods

The following methods are available on the `OutParameters` interface:

- `Object get(int)` (inherited from `Readable`)

- `Object get(String)` (inherited from `Readable`)

- `<T> T get(int, Class<T>)` (inherited from `Readable`)

- `<T> T get(String, Class<T>)` (inherited from `Readable`)

- `OutParametersMetadata getMetadata()`

`get(int[, Class])` methods accept parameter indexes starting at `0`, `get(String[, Class])` methods accept parameter names as they are represented in the result. Parameter names used as input to the `get` methods are case-insensitive. The following example shows how to create and consume `OutParameters` by using its index:

Example 35. Creating and Consuming `OutParameters` by obtaining a parameter by index

```
// result is a Result object
Publisher<Object> values = result.map((readable) -> readable.get(0));
```

The following example shows how to create and consume `OutParameters` by using a parameter name:

Example 36. Creating and Consuming a `OutParameters` by obtaining a parameter by name

```
// result is a Result object
Publisher<Object> titles = result.map((readable) -> readable.get("title"));
```

Calling `get` without specifying a target type returns a suitable value representation according to Mapping of Data Types. When specifying a target type, the R2DBC driver tries to convert the value to the target type. The following example shows how to create and consume `OutParameters` with type conversion:

Example 37. Creating and Consuming `OutParameters` *with type conversion*

```
// result is a Result object
Publisher<String> values = result.map((readable) -> readable.get(0,
String.class));
```

You can consume multiple parameters from `OutParameters`, as the following example shows:

Example 38. Consuming multiple parameters from `OutParameters`

```
// result is a Result object
Publisher<Book> values = result.map((readable) -> {
    String title = readable.get("title", String.class);
    String author = readable.get("author", String.class);

    return new Book(title, author);
});
```

When the parameter value in the database is SQL `NULL`, it can be returned to the Java application as `null`.

> ℹ️ `null` values cannot be returned as Reactive Streams values and must be wrapped for subsequent usage.

> ℹ️ Invalidating `OutParameters` does **not** release `Blob` and `Clob` objects that were obtained from `OutParameters`. These objects remain valid for at least the duration of the transaction in which they were created, unless their `discard()` method is called.

Chapter 11. Column and Row Metadata

The `RowMetadata` interface is implemented by R2DBC drivers to provide information about tabular results. It is used primarily by libraries and applications to determine the properties of a row and its columns.

In cases where the result properties of an SQL statement are unknown until it is run, the `RowMetadata` can be used to determine the actual properties of a row.

`RowMetadata` exposes `ColumnMetadata` for each column in the result. Drivers should provide `ColumnMetadata` on a best-effort basis. Column metadata is typically a by-product of statement execution. The amount of available information is vendor-dependent. Metadata retrieval can require additional lookups (internal queries) to provide a complete metadata descriptor. Issuing queries during result processing conflicts with the streaming nature of R2DBC. Consequently, `ColumnMetadata` declares two sets of methods: methods that must be implemented and methods that can optionally be implemented by drivers.

11.1. Obtaining a `RowMetadata` Object

A `RowMetadata` object is created during tabular results consumption through `Result.map(⋯)`. It is created for each row.The following example illustrates retrieval and usage by using an anonymous inner class:

Example 39. Using `RowMetadata` and retrieving `ColumnMetadata`

```
// result is a Result object
result.map(new BiFunction<Row, RowMetadata, Object>() {

    @Override
    public Object apply(Row row, RowMetadata rowMetadata) {
        ColumnMetadata my_column = rowMetadata.getColumnMetadata("my_column");
        Nullability nullability = my_column.getNullability();
        // ⋯
    }
});
```

11.2. Retrieving `ColumnMetadata`

`RowMetadata` methods are used to retrieve metadata for a single column or all columns.

- `getColumnMetadata(int)` returns the `ColumnMetadata` by using a zero-based index. See [Guidelines and Requirements](#).

- `getColumnMetadata(String)` returns the `ColumnMetadata` by using the column name (or alias as it is represented in the result).

- `getColumnMetadatas()` returns an unmodifiable `List` of `ColumnMetadata` objects.

- `contains(String)` returns whether `RowMetadata` contains metadata for the given column name. The column name (or alias as it is represented in the result) uses case-insensitive comparison rules.

11.3. Retrieving General Information for a Column

`ColumnMetadata` declares methods to access column metadata on a best-effort basis. Column metadata that is available as a by-product of running a statement must be made available through `ColumnMetadata`. Metadata exposure requiring interaction with the database (for example, issuing queries to information schema entities to resolve type properties) cannot be exposed, because methods on `ColumnMetadata` are expected to be non-blocking.

> ℹ Implementation note: Drivers can use metadata from a static mapping or obtain metadata indexes on connection creation.

The following example shows how to consume `ColumnMetadata` by using lambdas:

Example 40. Retrieving `ColumnMetadata` *information*

```
// row is a RowMetadata object
row.getColumnMetadatas().forEach(columnMetadata -> {

    String name = columnMetadata.getName();
    Integer precision = columnMetadata.getPrecision();
    Integer scale = columnMetadata.getScale();
});
```

See the API specification for more details.

Chapter 12. OUT Parameter Metadata

The `OutParametersMetadata` interface is implemented by R2DBC drivers to provide information about OUT parameters. It is used primarily by libraries and applications to determine the properties of result parameters.

In cases where the result properties of an SQL statement are unknown until it is run, the `OutParametersMetadata` can be used to determine the actual properties.

`OutParametersMetadata` exposes `OutParameterMetadata` for each returned out parameter. Drivers should provide `OutParameterMetadata` on a best-effort basis. Out parameter metadata is typically a by-product of stored procedure execution. The amount of available information is vendor-dependent. Metadata retrieval can require additional lookups (internal queries) to provide a complete metadata descriptor. Issuing queries during result processing conflicts with the streaming nature of R2DBC. Consequently, `OutParameterMetadata` declares two sets of methods: methods that must be implemented and methods that can optionally be implemented by drivers.

12.1. Obtaining a `OutParametersMetadata` Object

A `OutParametersMetadata` object is created during results consumption through `Result.map(⋯)`. It is created once for all out parameters. The following example illustrates retrieval and usage by using an anonymous inner class:

Example 41. Using `OutParametersMetadata` and retrieving `OutParameterMetadata`

```java
// result is a Result object
result.map(new Function<Readable, Object>() {

    @Override
    public Object apply(Readable readable) {

        if (readable instanceof OutParameters) {

            OutParameters out = (OutParameters) readable;
            OutParametersMetadata md = out.getMetadata();

            OutParameterMetadata my_parameter =
md.getParameterMetadata("my_parameter");
            Nullability nullability = my_parameter.getNullability();
            // ⋯
        }
    }
});
```

12.2. Retrieving `OutParameterMetadata`

`OutParameterMetadata` methods are used to retrieve metadata for a single parameter or all parameters.

- `getParameterMetadata(int)` returns the `OutParameterMetadata` by using a zero-based index. See Guidelines and Requirements.

- `getParameterMetadata(String)` returns the `OutParameterMetadata` by using the parameter name.

- `getParameterMetadatas()` returns an unmodifiable `List` of `OutParameterMetadata` objects.

12.3. Retrieving General Information for a OUT Parameter

`OutParameterMetadata` declares methods to access out parameter metadata on a best-effort basis. Out parameter metadata that is available as a by-product of running a statement must be made available through `OutParameterMetadata`. Metadata exposure requiring interaction with the database (for example, issuing queries to information schema entities to resolve type properties) cannot be exposed, because methods on `OutParameterMetadata` are expected to be non-blocking.

> ℹ️ Implementation note: Drivers can use metadata from a static mapping or obtain metadata indexes on connection creation.

The following example shows how to consume `OutParameterMetadata` by using lambdas:

Example 42. Retrieving `OutParameterMetadata` information

```
// out is a OutParametersMetadata object
out.getParameterMetadatas().forEach(outParameterMetadata -> {

    String name = outParameterMetadata.getName();
    Integer precision = outParameterMetadata.getPrecision();
    Integer scale = outParameterMetadata.getScale();
});
```

See the API specification for more details.

Chapter 13. Exceptions

This section explains how R2DBC uses and declares exceptions to provide information about various types of failures.

An exception is thrown by a driver when an error occurs during interaction with the driver or a data source. R2DBC differentiates between generic and data-source-specific error cases.

13.1. General Exceptions

R2DBC defines the following general exceptions:

- `IllegalArgumentException`
- `IllegalStateException`
- `UnsupportedOperationException`
- `NoSuchOptionException`
- `R2dbcException`

13.1.1. `IllegalArgumentException`

Drivers throw `IllegalArgumentException` if a method has received an illegal or inappropriate argument (such as values that are invalid or an expected parameter is `null`). This exception is a generic exception that is not associated with an error code or an `SQLState`.

13.1.2. `IndexOutOfBoundsException`

Drivers throw `IndexOutOfBoundsException` if a method has received an illegal or inappropriate index value (such as negative values or indexes that are greater or equal to the number of elements). This exception is a generic exception that is not associated with an error code or an `SQLState`.

13.1.3. `IllegalStateException`

Drivers throw `IllegalStateException` if a method has received an argument that is invalid in the current state or when an argument-less method is invoked in a state that does not allow execution in the current state (such as interacting with a closed connection object). This exception is a generic exception that is not associated with an error code or an `SQLState`.

13.1.4. `NoSuchElementException`

Drivers throw `NoSuchElementException` if a method has received a column/parameter name that doesn't exist in the collection of items. This exception is a generic exception that is not associated with an error code or an `SQLState`.

13.1.5. `NoSuchOptionException`

R2DBC SPI throws `NoSuchOptionException` if a required `Option` is retrieved from

`ConnectionFactoryOptions` that is either not configured or not associated with a value. `NoSuchOptionException` is a subclass of `IllegalStateException`.

13.1.6. `UnsupportedOperationException`

Drivers throw `UnsupportedOperationException` if the driver does not support certain functionality (such as when a method implementation cannot be provided). This exception is a generic exception that is not associated with an error code or an `SQLState`.

13.1.7. `R2dbcException`

Drivers throw an instance of `R2dbcException` when an error occurs during an interaction with a data source.

The exception contains the following information:

- A textual description of the error. You can retrieve the `String` that contains the description by invoking `R2dbcException.getMessage()`. Drivers may provide a localized message variant.

- An `SQLState`. The `String` that contains the `SQLState` can be retrieved by calling the `R2dbcException.getSqlState()` method. The value of the `SQLState` string depends on the underlying data source.

- An error code. The code is an integer value that identifies the error that caused the `R2dbcException` to be thrown. Its value and meaning are implementation-specific and may be the actual error code returned by the underlying data source. You can retrieve the error code by using the `R2dbcException.getErrorCode()` method.

- The offending SQL. If a particular SQL statement results in an exception then you can can retrieve the SQL text using the `R2dbcException.getSql()` method.

- A cause. This is another `Throwable` that caused this `R2dbcException` to occur.

13.2. Categorized Exceptions

Categorized exceptions provide a standard mapping to common error states. An R2DBC driver should provide specific subclasses to indicate affinity with the driver. Categorized exceptions provide a standardized approach for R2DBC clients and R2DBC users to translate common exceptions into an application-specific state without the need to implement `SQLState`-based exception translation, resulting in more portable error-handling code.

R2DBC categorizes exceptions into two top-level categories:

- Non-Transient Exceptions
- Transient Exceptions

13.2.1. Non-Transient Exceptions

A non-transient exception must extend the abstract class, `R2dbcNonTransientException`. A non-transient exception is thrown when a retry of the same operation would fail unless the cause of the error is corrected. After a non-transient exception other than `R2dbcNonTransientResourceException`,

the application can assume that a connection is still valid.

R2DBC defines the following subclasses of non-transient exceptions:

- `R2dbcBadGrammarException`: Thrown when the SQL statement has a problem in its syntax.
- `R2dbcDataIntegrityViolationException`: Thrown when an attempt to insert or update data results in a violation of an integrity constraint.
- `R2dbcPermissionDeniedException`: Thrown when the underlying resource denied a permission to access a specific element, such as a specific database table.
- `R2dbcNonTransientException`: Thrown when a resource fails completely and the failure is permanent. A connection may not be considered valid if this exception is thrown.

13.2.2. Transient Exceptions

A transient exception must extend the abstract class, `R2dbcTransientException`. A transient exception is thrown when a previously failed operation might be able to succeed if the operation is retried without any intervention in application-level functionality. After a non-transient exception other than `R2dbcTransientResourceException`, the application may assume that a connection is still valid.

- `R2dbcRollbackException`: Thrown when an attempt to commit a transaction resulted in an unexpected rollback due to deadlock or transaction serialization failures.
- `R2dbcTimeoutException`: Thrown when the timeout specified by a database operation (query, login, and so on) is exceeded. This could have different causes (depending on the database API in use) but is most likely thrown after the database interrupts or stops the processing of a query before it has completed.
- `R2dbcNonTransientException`: Thrown when a resource fails temporarily and the operation can be retried. A connection may not be considered valid if this exception is thrown.

Chapter 14. Data Types

This chapter discusses the use of data types from Java and database perspectives. The R2DBC SPI gives applications access to data types that are defined as SQL. R2DBC is not limited to SQL types, and, in fact, the SPI is type-agnostic.

If a data source does not support a data type described in this chapter, a driver for that data source is not required to implement the methods and interfaces associated with that data type.

14.1. Mapping of Data Types

This section explains how SQL-specific types map to Java types. The list is not exhaustive and should be received as a guideline for drivers. R2DBC drivers should use modern types and type descriptors to exchange data for consumption by applications and consumption by the driver. Driver implementations should implement the following type mapping and can support additional type mappings:

- Character Types
- Boolean Types
- Binary Types
- Numeric Types
- Datetime Types
- Collection Types

The following table describes the SQL type mapping for character types:

Table 4. SQL Type Mapping for Character Types

SQL Type	Description	Java Type
`CHARACTER` (`CHAR`)	Character string, fixed length.	`java.lang.String`
`CHARACTER VARYING` (`VARCHAR`)	Variable-length character string, maximum length fixed.	`java.lang.String`
`NATIONAL CHARACTER` (`NCHAR`)	The `NATIONAL CHARACTER` type is the same as `CHARACTER` except that it holds standardized multibyte characters or Unicode characters.	`java.lang.String`
`NATIONAL CHARACTER VARYING` (`NVARCHAR`)	The `NATIONAL CHARACTER VARYING` type is the same as `CHARACTER VARYING` except that it holds standardized multibyte characters or Unicode characters.	`java.lang.String`

SQL Type	Description	Java Type
CHARACTER LARGE OBJECT (CLOB)	A Character Large OBject (or CLOB) is a collection of character data in a DBMS, usually stored in a separate location that is referenced in the table itself. Note that drivers may default to Clob when materializing a CLOB value requires additional database communication.	java.lang.String, io.r2dbc.spi.Clob
NATIONAL CHARACTER LARGE OBJECT (NCLOB)	The NATIONAL CHARACTER LARGE OBJECT type is the same as CHARACTER LARGE OBJECT except that it holds standardized multibyte characters or Unicode characters. Note that drivers may default to Clob when materializing a NCLOB value requires additional database communication.	java.lang.String, io.r2dbc.spi.Clob

The following table describes the SQL type mapping for boolean types:

Table 5. SQL Type Mapping for Boolean Types

SQL Type	Description	Java Type
BOOLEAN	A value that represents a boolean state.	java.lang.Boolean

The following table describes the SQL type mapping for binary types:

Table 6. SQL Type Mapping for Binary Types

SQL Type	Description	Java Type
BINARY	Binary data, fixed length.	java.nio.ByteBuffer
BINARY VARYING (VARBINARY)	A variable-length character string, the maximum length of which is fixed.	java.nio.ByteBuffer

SQL Type	Description	Java Type
BINARY LARGE OBJECT (BLOB)	A Binary Large OBject (or BLOB) is a collection of binary data in a database management system, usually stored in a separate location that is referenced in the table itself. Note that drivers may default to Blob when materializing a BLOB value requires additional database communication.	java.nio.ByteBuffer, io.r2dbc.spi.Blob

The following table describes the SQL type mapping for numeric types:

Table 7. SQL Type Mapping for Numeric Types

SQL Type	Description	Java Type
INTEGER	Represents an integer. The minimum and maximum values depend on the DBMS (typically 4-byte precision).	java.lang.Integer
TINYINT	Same as the INTEGER type except that it might hold a smaller range of values, depending on the DBMS (typically 1-byte precision).	java.lang.Byte
SMALLINT	Same as the INTEGER type except that it might hold a smaller range of values, depending on the DBMS (typically 1- or 2-byte precision).	java.lang.Short
BIGINT	Same as the INTEGER type except that it might hold a larger range of values, depending on the DBMS (typically 8-byte precision).	java.lang.Long
DECIMAL(p, s), NUMERIC(p, s)	Fixed precision and scale numbers with precision (p) and scale (s). In other words, a number that can have a decimal point in it.	java.math.BigDecimal

SQL Type	Description	Java Type
`FLOAT(p)`	Represents an approximate numerical with mantissa precision (`p`). Databases that use IEEE representation can map values to either 32-bit or 64-bit floating point types depending on the precision parameter (`p`).	`java.lang.Double` or `java.lang.Float`
`REAL`	Same as the `FLOAT` type except that the DBMS defines the precision.	`java.lang.Float`
`DOUBLE PRECISION`	Same as the `FLOAT` type except that the DBMS defines the precision. It has greater precision than `REAL`.	`java.lang.Double`

The following table describes the SQL type mapping for datetime types:

Table 8. SQL Type Mapping for Datetime Types

SQL Type	Description	Java Type
`DATE`	Represents a date without specifying a time part and without a timezone.	`java.time.LocalDate`
`TIME`	Represents a time without a date part and without a timezone.	`java.time.LocalTime`
`TIME WITH TIME ZONE`	Represents a time with a timezone offset.	`java.time.OffsetTime`
`TIMESTAMP`	Represents a date and time without a timezone.	`java.time.LocalDateTime`
`TIMESTAMP WITH TIME ZONE`	Represents a date and time with a timezone offset.	`java.time.OffsetDateTime`

The following table describes the SQL type mapping for collection types:

Table 9. SQL Type Mapping for Collection Types

SQL Type	Description	Java Type
`COLLECTION` (`ARRAY`, `MULTISET`)	Represents a collection of items with a base type.	Array-Variant of the corresponding Java type (for example, `Integer[]` for `INTEGER ARRAY`)

Vendor-specific types (such as spatial data types, structured JSON or XML data, and user-defined types) are subject to vendor-specific mapping.

14.2. Type Descriptors

R2DBC drivers may infer the database type for inbound parameters or use a specific type. R2DBC's type system `io.r2dbc.spi.Type` and `io.r2dbc.spi.Parameter` are interfaces to describe a database type and a typed parameter. The R2DBC specification defines its type mapping in the `io.r2dbc.spi.R2dbcType` utility for commonly used data types. R2DBC drivers may provide their own `Type` objects to provide vendor-specific type support.

14.3. Mapping of Advanced Data Types

The R2DBC SPI declares default mappings for advanced data types.The following list describes data types and the interfaces to which they map:

- `BLOB`: The `Blob` interface
- `CLOB`: The `Clob` interface

14.3.1. `Blob` and `Clob` Objects

An implementation of a `Blob` or `Clob` object may either be locator-based or fully materialize the object in the driver. Drivers should prefer locator-based `Blob` and `Clob` interface implementations to reduce pressure on the client when materializing results.

For implementations that fully materialize Large OBjects (LOBs), the `Blob` and `Clob` objects remain valid until the LOB is consumed or the `discard()` method is called.

Portable applications should not depend upon the LOB validity past the end of a transaction.

14.3.2. Creating `Blob` and `Clob` Objects

Large objects are backed by a `Publisher` that emits the component type of the large object, such as `ByteBuffer` for `BLOB` and `CharSequence` (or a subtype of it) for `CLOB`.

Both interfaces provide factory methods to create implementations to be used with `Statement`. The following example shows how to create a `Clob` object:

Example 43. Creating and using a `Clob` object

```
// charstream is a Publisher<String> object
  // statement is a Statement object
Clob clob = Clob.from(charstream);
statement.bind("text", clob);
```

14.3.3. Retrieving `Blob` and `Clob` Objects from a `Readable`

The Binary Large OBject (`BLOB`) and Character Large OBject (`CLOB`) data types are treated similarly to primitive built-in types. You can retrieve values of these types by calling the `get(…)` methods on the

`Readable` interface. The following example shows how to do so:

Example 44. Retrieving a `Clob` object

```
// result is a Result object
Publisher<Clob> clob = result.map((readable) -> readable.get("clob", Clob.class));
```

The `Blob` and `Clob` interfaces contain methods for returning the content and for releasing resources associated with the object instance. The API documentation provides more details.

> ℹ️ LOB value consumption requires special attention due to large object capacity needs. Mapping of Data Types defines a default LOB mapping using scalar types. The actual default data type for LOB data types can be vendor-specific to avoid blocking if the database requires LOB materialization from a locator or requires database communication during retrieval.

14.3.4. Accessing `Blob` and `Clob` Data

The `Blob` and `Clob` interfaces declare methods to consume the content of each type. Content streams follow Reactive Streams specifications and reflect the stream nature of large objects. As a result, `Blob` and `Clob` objects can be consumed only once. Large object data consumption can be canceled by calling the `discard()` method if the content stream was not consumed at all.Alternatively, if the content stream was consumed, a `Subscription` cancellation releases resources that are associated with the large object.

The following example shows how to consume `Clob` contents:

Example 45. Creating and using a `Clob` object

```
// clob is a Clob object
Publisher<CharSequence> charstream = clob.stream();
```

14.3.5. Releasing `Blob` and `Clob`

`Blob` and `Clob` objects remain valid for at least the duration of the transaction in which they are created. This could potentially result in an application running out of resources during a long-running transaction. Applications may release `Blob` and `Clob` by either consuming the content stream or disposing of resources by calling the `discard()` method.

The following example shows how to free `Clob` resources without consuming it:

Example 46. Freeing `Clob` *object resources*

```
// clob is a Clob object
Publisher<Void> charstream = clob.discard();
charstream.subscribe(…);
```

Chapter 15. Extensions

This section covers optional extensions to the [R2DBC Core](#). Extensions provide features that are not mandatory for R2DBC implementations.

15.1. Wrapped Interface

The `Wrapped` interface provides a way to access an instance of a resource which has been wrapped and for implementors to expose wrapped resources. This mechanism helps to eliminate the need to use non-standard means to access vendor-specific resources.

15.1.1. Usage

A wrapper for a R2DBC SPI type is expected to implement the `Wrapped` interface so that callers can extract the original instance. Any R2DBC SPI interface type can be wrapped. The following example shows how to expose a wrapped resource:

Example 47. Wrapping a `Connection` and exposing the underlying resource.

```
class ConnectionWrapper implements Connection, Wrapped<Connection> {

    private final Connection wrapped;

    @Override
    public Connection unwrap() {
        return this.wrapped;
    }

    // constructors and implementation methods omitted for brevity.
}
```

15.1.2. Interface Methods

The following methods are available on the `Wrapped` interface:

* `unwrap`

15.1.3. The `unwrap` Method

The `unwrap` method is used to return an object that implements the specified interface, allowing access to vendor-specific methods. The returned object may either be the object found to implement the specified interface or a wrapper for that object. Wrappers can be unwrapped recursively. The following example shows how to unwrap a wrapped object:

Example 48. Unwrapping a wrapped object.

```
// connection is a Connection object implementing Wrapped

if (connection instanceof Wrapped) {
    connection = ((Wrapped<Connection>) connection).unwrap();
}
```

15.2. Closeable Interface

The `io.r2dbc.spi.Closeable` interface provides a mechanism for objects associated with resources to release these resources once the object is no longer in use. The associated resources are released without blocking the caller.

15.2.1. Usage

A closeable object is expected to implement the `Closeable` interface so that callers can obtain a `Publisher` to initiate the close operation and get notified upon completion. The following example shows how to close a connection:

Example 49. Closing a `Connection`.

```
// connection is a Connection object
Publisher<Void> close = connection.close();
```

`Connection` implements `Closeable` as a mandatory part of R2DBC. Any stateful object (such as `ConnectionFactory`) can implement `Closeable` to provide a way to release its resources.

15.2.2. Interface Methods

The following methods are available on the `Closeable` interface:

- `close`

15.2.3. The `close` Method

The `close` method is used to return a `Publisher` to start the close operation and get notified upon completion. If the object is already closed, then subscriptions complete successfully and the close operation has no effect.

15.3. Lifecycle Interface

The `Lifecycle` interface provides methods to notify a connection resource about its lifecycle state. Typically used by resource management components such as connection pools to indicate

allocation and release phases so that connections may allocate or release resources as needed.

15.3.1. Usage

The `Lifecycle` interface is typically implemented by `Connection` objects. Those are notified by their resource manager such as a pool right before returning a cached connection for usage and after usage, right before returning the connection to the pool or closing the connection. Connections implementing `Lifecycle` may allocate resources upon `postAllocate` and release these before going into idle state on `preRelease`. Lifecycle methods can be used to clean up unfinished transactions or reset the connection state. Over the lifespan of a cached resource, lifecycle methods are called multiple times in the sequence of:

1. `ConnectionFactory.create()`
2. `postAllocate`
3. Application-specific work
4. `preRelease`
5. Repeat from 2. or `Connection.close()`

Error signals emitted by lifecycle method publishers should be propagated appropriately to the caller.

15.3.2. Interface Methods

The following methods are available on the `Lifecycle` interface:

- `postAllocate`
- `preRelease`

15.3.3. The `postAllocate` Method

The `postAllocate` method is used to return a `Publisher` to signal allocation to a resource. Any application-specific work happens after completion of the `Publisher` returned by this method. Successive calls to `postAllocate` are expected to no-op.

15.3.4. The `preRelease` Method

The `preRelease` method is used to return a `Publisher` to signal release of a resource to idle state. Any application-specific work has completed before subscribing to the `Publisher` returned by this method. Successive calls to `preRelease` are expected to no-op.